

A04 - Assignment 4

- Deidier Simone - *student number: 133020*

Theory questions

1. In the shopping cart example, **version numbers** capture the *happens-before* relationship by marking the sequence of updates, in fact when two concurrent updates to the cart are made (e.g., from different devices), version numbers act as **logical clocks**, recording the time of each action. These numbers help detect conflicts (i.e., concurrent changes needing a merge), ensuring that each device has a **consistent cart state** by deciding the most recent or conflict-resolved version.
2. According to Kleppmann, the best re-partitioning method is *consistent hashing*. This approach **minimizes the number of keys that need to move** when partition count changes, reducing re-partitioning overhead. *Consistent hashing* also provides scalable load balancing as nodes or partitions are added or removed.
 - **Local indexing** is ideal for databases where each partition has its index, enabling **faster, independent queries** on each shard, but it's less suitable for **cross-partition queries**.
 - **Global indexing** is needed for queries covering multiple partitions, providing a **global view**. While it's more complex and slower, it ensures *data consistency* across partitions.
3. Read Committed vs. Snapshot Isolation for the schedule $r_1(A); w_2(A); w_2(B); r_1(B); c_1; c_2$:
 - **Read Committed:** T1 reads A before T2 writes it, seeing the original value. T1 reads B after T2 writes it, seeing the **new value of B**. **Outcome:** T1 might see *mixed old and new values*, leading to potential anomalies.
 - **Snapshot Isolation:** T1 takes a *consistent snapshot* of A and B from the start. **Outcome:** T1 sees the same versions of A and B, avoiding anomalies present in *Read Committed*.
 - With **2PL**, T1 and T2 adhere to the two-phase locking protocol: **locks prevent conflicting reads and writes** until transactions complete, this approach guarantees serializability by **preventing anomalies** from concurrent access.
4. Possible reasons for no reply to a network message are **network latency** or congestion, the receiver node may have *crashed or become unreachable*, the message may be *lost in transit* or dropped, acknowledgment failure or *firewall/security* rule blocking message return.
 - Using clocks for *last write wins* is dangerous due to **clock drift** between nodes, which can cause inaccurate conflict resolution. One

node's "newer" write may appear "older" on another, leading to inconsistencies.

- A node may not trust its judgment in cases like **network partitions** or **clock drift**, where it might believe it has the latest data but cannot confirm with other nodes.
5. **Ordering** ensures operations apply in a set sequence, **linearizability** guarantees a strong *consistency model*, where each operation appears to occur at a specific instant, and **consensus** (e.g., RAFT, Paxos) ensures all nodes agree on an operation order, achieving *linearizability* across a distributed system.
- *Non-linearizable systems* like eventual consistency databases are widely used where strict consistency isn't necessary, in fact they offer high availability and partition tolerance, fitting **large-scale applications**.
 - Ensuring linearizability can reduce **system performance** and increase **latency**. The high cost of coordination often makes linearizability unsuitable for high-performance, highly available systems.
6. With **logical clocks**, $L(e) < L(f)$ doesn't guarantee e happened before f due to possible concurrency. However, *vector clocks* can capture causality, so if $V(e) < V(f)$, then e indeed happened before f . Vector clock values for remaining events in the figure:

- $b \rightarrow (4, 0, 0)$
- $k \rightarrow (4, 2, 0)$
- $m \rightarrow (4, 3, 0)$
- $c \rightarrow (4, 3, 2)$
- $u \rightarrow (4, 4, 0)$
- $n \rightarrow (5, 4, 0)$

Alternative consistent states from state S_{00} :

- S_{00} : (0, 0) for P_1 and (0, 0) for P_2 (initial state)
 - S_{10} : (1, 0) for P_1 , (0, 0) for P_2 (P_1 has completed an event, but no messages exchanged yet)
 - S_{11} : (1, 0) for P_1 , (0, 1) for P_2 (P_2 has advanced independently of P_1)
 - S_{20} : (2, 0) for P_1 , (0, 1) for P_2 (P_1 advances again without communication)
 - S_{21} : (2, 0) for P_1 , (2, 1) for P_2 (P_2 receives the message from P_1 , updating its clock)
 - S_{22} : (2, 2) for P_1 , (2, 2) for P_2 (P_2 performs an event and sends a message to P_1)
 - S_{33} : (3, 3) for P_1 , (2, 3) for P_2 (P_1 receives the message from P_2 , and both clocks are synchronized)
7. RAFT ensures **log consistency** by requiring new leaders to synchronize logs with the majority before appending new entries. The leader **replicates its log to followers**, overwriting inconsistencies to maintain uniform logs across all nodes.

8. The main advantages of RAFT in MySQL are **improved consensus and fault tolerance, consistency across replicas** and **simplifies leader election and log replication**, enhancing MySQL's reliability for distributed transactions.
9. Stonebraker and Pavlo note that while **document databases** (e.g., MongoDB) are flexible for unstructured data, *SQL databases* have regained popularity for their **ACID guarantees** and **complex querying power**. A convergence trend shows SQL databases adopting schema flexibility, while NoSQL systems strengthen transaction guarantees, blending benefits across both types.